

Operating system

Lab 5



Author: 吴杭 李克成

Date: 2024/12/24

Autumn-Fall 2024-2025 Semester

Table of Contents

Chapter 1: 实验流程	3
1.1) 4.3 实现 fork 的系统调用	3
1.1.1) 拷贝内核栈	3
1.1.2) 创建子进程页表	3
1.1.3) 处理进程返回逻辑	5
1.2) 运行结果	5
Chapter 2: 思考题	11
Chapter 3: 心得体会	11
3.1) 遇到的问题	11
3.2) 心得体会	12
Declaration	12

Chapter 1: 实验流程

1.1) 4.3 实现 fork 的系统调用

首先根据实验文档中的内容搭好框架，这里不再展开。

1.1.1) 拷贝内核栈

直接根据 current（当前在运行的线程的 task struct）来深拷贝这一页的内容，就能把内核栈拷贝到新的进程中。其他再根据实验文档来修改 task_struct 相关的内容。

```
uint64_t do_fork(struct pt_regs *regs) {
    // copy kernel stack
    struct task_struct *_task = (struct task_struct *)kalloc();
    memcpy(_task, current, PGSIZE);
    // assign new pid
    _task->pid = nr_tasks++;
    // assign new page directory
    _task->pgd = alloc_page();
    // set mm.mmap
    _task->mm.mmap = NULL;
    ...
}
```

1.1.2) 创建子进程页表

内核页表每个线程都是一致的，所以直接拷贝就行。

```
// copy the swapper_pg_dir into the new task's pgd
memcpy(_task->pgd, swapper_pg_dir, PGSIZE);
```

接下来需要遍历父进程的所有 vma，子进程作为父进程的副本，此时的运行状态理论上要和父进程一致，所以需要把父进程的所有 vma 都复制一份，添加到子进程的 vma 链表中。同时，父进程的 vma 对应的页表如果存在，说明此时这一页已经被创建，所以需要把这一页的内容拷贝到子进程的页表中。

下面是拷贝 vma 的代码，主要的逻辑就是复制一份 vma 的结构，然后执行链表的 insert 操作即可。

```
// iterate the vma list of the current task
struct vm_area_struct *vma = current->mm.mmap;
while (vma != NULL) {
    // create a new vma
    struct vm_area_struct *new_vma = (struct vm_area_struct *)kalloc();
    memcpy(new_vma, vma, sizeof(struct vm_area_struct));
    // link the new vma to the new task's mm
    new_vma->vm_mm = &(_task->mm);
    new_vma->vm_next = _task->mm.mmap;
    new_vma->vm_prev = NULL;
    // link the new vma to the new task's mm
    if (_task->mm.mmap == NULL) {
        _task->mm.mmap = new_vma;
    } else {
```

```

        // struct vm_area_struct *tmp = _task->mm.mmap;
        _task->mm.mmap->vm_prev = new_vma;
        _task->mm.mmap = new_vma;
    }
    // if the corresponding page table entry exists, copy the content
    check_and_copy_pages(current->pgd, new_vma->vm_start, new_vma->vm_end, _task-
>pgd, new_vma->vm_flags);
    // move to the next vma
    vma = vma->vm_next;
}

```

下面是拷贝 vma 对应的页表项的代码，主要的逻辑是通过传入的 va_start 和 va_end 来遍历经过的每一页，然后通过页表项的 vpn 来找到对应的页表项，然后判断对应页表项的 PRIV_V，如果存在，就给予进程创建一个新的页，再通过 vma 记录的页权限来 creat_mapping，并且把父进程的页深拷贝一份到子进程的页中。这样就保证了父进程和子进程现在用到的资源是完全一致的。

```

// utils: check and copy the content of the page table entry
void check_and_copy_pages(uint64_t *pgd, uint64_t va_start, uint64_t va_end, uint64_t
*new_pgd, uint64_t vm_flags) {
    // notice va is page aligned
    for (uint64_t va = PGROUNDDOWN(va_start); va < va_end; va += PGSIZE) {
        // get the page table entry
        uint64_t vpn0 = (va >> 12) & 0x1ff;
        uint64_t vpn1 = (va >> 21) & 0x1ff;
        uint64_t vpn2 = (va >> 30) & 0x1ff;

        // check if the first page table entry is valid
        if (pgd[vpn2] & PRIV_V) {
            uint64_t *pgtbl1 = get_pgtable(pgd, vpn2);
            // check if the second page table entry is valid
            if (pgtbl1[vpn1] & PRIV_V) {
                uint64_t *pgtbl0 = get_pgtable(pgtbl1, vpn1);
                // check if the third page table entry is valid
                if (pgtbl0[vpn0] & PRIV_V) {
                    // if yes, deep copy the content of the page
                    // notice only the machine mode has the privilege to access the
physical address, so we use memcpy for the virtual page address here
                    char* child_process_page = alloc_page();
                    uint64_t priv_r = (vm_flags & VM_READ) ? PRIV_R : 0;
                    uint64_t priv_w = (vm_flags & VM_WRITE) ? PRIV_W : 0;
                    uint64_t priv_x = (vm_flags & VM_EXEC) ? PRIV_X : 0;
                    create_mapping(new_pgd, va, VA2PA((uint64_t)child_process_page),
PGSIZE, PRIV_U | priv_r | priv_w | priv_x | PRIV_V);
                    // copy the content of current page to the child page
                    memcpy(child_process_page, (char *)va, PGSIZE);
                }
            }
        }
    }
}

```

1.1.3) 处理进程返回逻辑

对于父进程的返回，文档里面已经写得很清楚了，在运行完 `do_fork` 之后就正常返回子进程的 `pid` 即可，我们无需做任何修改。

而子进程可以认为是下一次被调度的时候才会返回，并且调度的过程中 `PC` 实际上是在 `__switch_to` 的位置，在 `__switch_to` 中，会把子进程的 `task_struct` 给 load 出来，然后根据其中设置的 `ra` 来跳转到下一个执行的地址。为了让子进程和父进程在 `fork` 时候的运行状态一致，我们应该让子进程像父进程那样返回到 `trap_handler` 返回的位置，也就是 `_traps` 中 `call trap_handler` 之后的位置加一个标记 `__ret_from_fork`，然后把子进程的 `thread.ra` 设置为这个位置即可。

接下来我们处理子进程的一系列 `sp`，首先是 `_task->thread.sp`，由于我们在 `__switch_to` 返回到 `__ret_from_fork` 的时候，马上需要根据此时的 `sp` 来 load 子进程存在内核栈中的值，而内核栈的位置可以通过 `do_fork` 中 `regs` 和 `current` 的相对位置以及 `_task` 的值来得到。

接下来是 `_task->thread.sscratch` 的值，理论上这个寄存器在此时应当存储子进程用户态栈的地址。而这个地址实际上保存在内核栈 `regs` 中存好的 `sscratch` 中。

另外我们还需要思考子进程的内核栈 `child_pt_regs` 中的 `sp` 的值。这个值根据我们前面的分析，实际上会在 `__ret_from_traps` 被 load 到 `sp` 中，观察 `_traps` 中的代码，就会发现他应当和 `_task->thread.sp` 的值相同，也就是和 `child_pt_regs` 的值相同。

最后我们设置子进程的返回值，也就是 `child_pt_regs` 中存好的 `a0`，将其设置为 0。以及设置子进程的 `sepc`，也存在 `sepc` 中，将其的值加 4 即可。

同时由于我们在 `task_struct` 中储存了 `satp`，所以也要作相应的更新。

```
// struct pt_regs* test_regs = (struct pt_regs *)((uint64_t)current + PGSIZE -
sizeof(struct pt_regs));
struct pt_regs* child_pt_regs = (struct pt_regs *)((uint64_t)_task + (uint64_t)regs -
(uint64_t)current);
child_pt_regs->general_regs[1] = (uint64_t)child_pt_regs; // set child_pt_regs->sp
child_pt_regs->sepc += 4; // set the sepc of the child process
_task->thread.sscratch = regs->sscratch; // the sscratch of the child process
_task->thread.sp = (uint64_t)child_pt_regs; // the sp of the child process
_task->thread.ra = (uint64_t)__ret_from_fork; // the ra of the child process
// set the return value of the child process
child_pt_regs->general_regs[9] = 0;
// set the satp of the child process
uint64_t ppn = VA2PA((uint64_t)(_task->pgd)) >> 12;
_task->thread.satp = ppn | (SATP_MODE_SV39 << 60);
// add _task to the task list
task[_task->pid] = _task;
// return the new task's pid
return _task->pid;
```

1.2) 运行结果

运行 `make run TEST=FORK1`

```

SET [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 12, sepc = 0x100e8
[trap.c,22,do_page_fault] page fault at 0x100e8
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802d1000
[trap.c,86,trap_handler] trap: scause = 15, sepc = 0x101ac
[trap.c,22,do_page_fault] page fault at 0x3fffffff8
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802de000
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 13, sepc = 0x10240
[trap.c,22,do_page_fault] page fault at 0x12000
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802e1000
[trap.c,86,trap_handler] trap: scause = 12, sepc = 0x1114c
[trap.c,22,do_page_fault] page fault at 0x1114c
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802e2000
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x1121c
[U-PARENT] pid: 1 is running! global_variable: 0
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 13, sepc = 0x101e8
[trap.c,22,do_page_fault] page fault at 0x12000
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802e3000
[trap.c,86,trap_handler] trap: scause = 12, sepc = 0x1114c
[trap.c,22,do_page_fault] page fault at 0x1114c
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802e4000
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x1121c
[U-CHILD] pid: 2 is running! global_variable: 0
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x1121c
[U-PARENT] pid: 1 is running! global_variable: 1
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x1121c
[U-CHILD] pid: 2 is running! global_variable: 1
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x1121c
[U-CHILD] pid: 2 is running! global_variable: 2
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100

```

Figure 1: “fork1”

运行 make run TEST=FORK2

```
2024 ZJU Operating System
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 12, sepc = 0x100e8
[trap.c,22,do_page_fault] page fault at 0x100e8
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802d2000, 0x802d3000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 15, sepc = 0x101ac
[trap.c,22,do_page_fault] page fault at 0x3fffffff8
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x4000000000) to [0x802d5000, 0x802d6000), perm = 0x17
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 13, sepc = 0x101d4
[trap.c,22,do_page_fault] page fault at 0x12000
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802d8000, 0x802d9000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 12, sepc = 0x11320
[trap.c,22,do_page_fault] page fault at 0x11320
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802d9000, 0x802da000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 13, sepc = 0x10470
[trap.c,22,do_page_fault] page fault at 0x14008
[vm.c,91,create_mapping] mapping [0x14000, 0x15000) to [0x802da000, 0x802db000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U] pid: 1 is running! global_variable: 0
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U] pid: 1 is running! global_variable: 1
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U] pid: 1 is running! global_variable: 2
[trap.c,86,trap_handler] trap: scause = 15, sepc = 0x10230
[trap.c,22,do_page_fault] page fault at 0x13008
[vm.c,91,create_mapping] mapping [0x13000, 0x14000) to [0x802db000, 0x802dc000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x4000000000) to [0x802df000, 0x802e0000), perm = 0x17
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802e3000, 0x802e4000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802e6000, 0x802e7000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802e7000, 0x802e8000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x13000, 0x14000) to [0x802e8000, 0x802e9000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x14000, 0x15000) to [0x802e9000, 0x802ea000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-PARENT] pid: 1 is running! Message: ZJU OS Lab5
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-PARENT] pid: 1 is running! global_variable: 3
switch to [PID = 2 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-CHILD] pid: 2 is running! Message: ZJU OS Lab5
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-CHILD] pid: 2 is running! global_variable: 3
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
```

Figure 2: “fork2”

```

[vm.c,91,create_mapping] mapping [0x14000, 0x15000) to [0x802e9000, 0x802ea000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-PARENT] pid: 1 is running! Message: ZJU OS Lab5
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-PARENT] pid: 1 is running! global_variable: 3
switch to [PID = 2 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-CHILD] pid: 2 is running! Message: ZJU OS Lab5
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-CHILD] pid: 2 is running! global_variable: 3
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-PARENT] pid: 1 is running! global_variable: 4
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-CHILD] pid: 2 is running! global_variable: 4
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-PARENT] pid: 1 is running! global_variable: 5
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-CHILD] pid: 2 is running! global_variable: 5
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-PARENT] pid: 1 is running! global_variable: 6
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x113f0
[U-CHILD] pid: 2 is running! global_variable: 6
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]

```

Figure 3: “fork2”

运行 make run TEST=FORK3


```

SET [PID = 1 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 12, sepc = 0x100e8
[trap.c,22,do_page_fault] page fault at 0x100e8
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802d1000, 0x802d2000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 15, sepc = 0x101ac
[trap.c,22,do_page_fault] page fault at 0x3fffffff8
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x802d4000, 0x802d5000), perm = 0x17
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 13, sepc = 0x101cc
[trap.c,22,do_page_fault] page fault at 0x12000
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802d7000, 0x802d8000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 12, sepc = 0x11174
[trap.c,22,do_page_fault] page fault at 0x11174
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802d8000, 0x802d9000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 1 is running! global_variable: 0
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x802dc000, 0x802dd000), perm = 0x17
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802e0000, 0x802e1000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802e3000, 0x802e4000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802e4000, 0x802e5000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x802e8000, 0x802e9000), perm = 0x17
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802ec000, 0x802ed000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802ef000, 0x802f0000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802f0000, 0x802f1000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 1 is running! global_variable: 1
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x802f4000, 0x802f5000), perm = 0x17
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x802f8000, 0x802f9000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x802fb000, 0x802fc000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x802fc000, 0x802fd000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 1 is running! global_variable: 2
switch to [PID = 2 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x80300000, 0x80301000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x80303000, 0x80304000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x80304000, 0x80305000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x80306000, 0x80307000), perm = 0x17
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 2 is running! global_variable: 1
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134

```

Figure 4: “fork3”

```

switch to [PID = 2 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x80300000, 0x80301000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x80303000, 0x80304000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x80304000, 0x80305000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x80306000, 0x80307000), perm = 0x17
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 2 is running! global_variable: 1
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x8030c000, 0x8030d000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x8030f000, 0x80310000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x80310000, 0x80311000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x80312000, 0x80313000), perm = 0x17
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 2 is running! global_variable: 2
switch to [PID = 3 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 3 is running! global_variable: 1
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x80318000, 0x80319000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x8031b000, 0x8031c000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x8031c000, 0x8031d000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x8031e000, 0x8031f000), perm = 0x17
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 3 is running! global_variable: 2
switch to [PID = 4 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 4 is running! global_variable: 2
switch to [PID = 5 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 5 is running! global_variable: 1
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10134
[trap.c,97,trap_handler] clone syscall
[vm.c,91,create_mapping] mapping [0x3fffffff000, 0x40000000000) to [0x80324000, 0x80325000), perm = 0x17
[vm.c,91,create_mapping] mapping [0x10000, 0x11000) to [0x80328000, 0x80329000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x11000, 0x12000) to [0x8032b000, 0x8032c000), perm = 0x1f
[vm.c,91,create_mapping] mapping [0x12000, 0x13000) to [0x8032c000, 0x8032d000), perm = 0x1f
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 5 is running! global_variable: 2
switch to [PID = 6 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244

```

Figure 5: “fork3”

```

[U] pid: 5 is running! global_variable: 2
switch to [PID = 6 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 6 is running! global_variable: 2
switch to [PID = 7 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 7 is running! global_variable: 2
switch to [PID = 8 PRIORITY = 7 COUNTER = 6]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 8 is running! global_variable: 2
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
SET [PID = 2 PRIORITY = 7 COUNTER = 7]
SET [PID = 3 PRIORITY = 7 COUNTER = 7]
SET [PID = 4 PRIORITY = 7 COUNTER = 7]
SET [PID = 5 PRIORITY = 7 COUNTER = 7]
SET [PID = 6 PRIORITY = 7 COUNTER = 7]
SET [PID = 7 PRIORITY = 7 COUNTER = 7]
SET [PID = 8 PRIORITY = 7 COUNTER = 7]
switch to [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 1 is running! global_variable: 3
switch to [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 2 is running! global_variable: 3
switch to [PID = 3 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 3 is running! global_variable: 3
switch to [PID = 4 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 4 is running! global_variable: 3
switch to [PID = 5 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 5 is running! global_variable: 3
switch to [PID = 6 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 6 is running! global_variable: 3
switch to [PID = 7 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244
[U] pid: 7 is running! global_variable: 3
switch to [PID = 8 PRIORITY = 7 COUNTER = 7]
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x10100
[trap.c,86,trap_handler] trap: scause = 8, sepc = 0x11244

```

Figure 6: “fork3”

Chapter 2: 思考题

Chpater 3: 心得体会

3.1) 遇到的问题

1. 在考虑子进程返回的 sp 设置的时候, 一开始认为_task->thread.spload 到 sp 中的值会在子进程进入到__ret_from_fork 后马上被内核栈中 load 出来的值给覆盖掉, 所以认

为 `_task->thread.sp` 不需要设置，后来才意识到在 `__ret_from_fork` 中需要是通过 `_task->thread.spload` 出来的 `sp` 的值来找到内核栈的位置，所以这个值是需要设置的。

3.2) 心得体会

这次实验让我对 `fork` 机制有了更加深刻的理解，尤其是对于 `fork` 之后，子进程和父进程返回值不同是怎么实现的，这一点之前困扰了我很久，在做了实验之后就很清晰了。这次 `do_fork` 中的程序虽然复杂，但是文档写的比较详细，按照上面的来也不会漏掉需要处理的结构。并且需要思考的部分也给了很清楚的提示，体验感很好。

Declaration

We hereby declare that all the work done in this lab 5 is of our independent effort.